


# NFT Price Prediction

**CS376 Machine Learning**

20170435 윤지언 20170653 차정엽 20170783 이동규 20190584 정서경

# Goal

Weekly total cryptocurrency value and average value per transaction sent to NFT platforms, 2021 - 2022 YTD

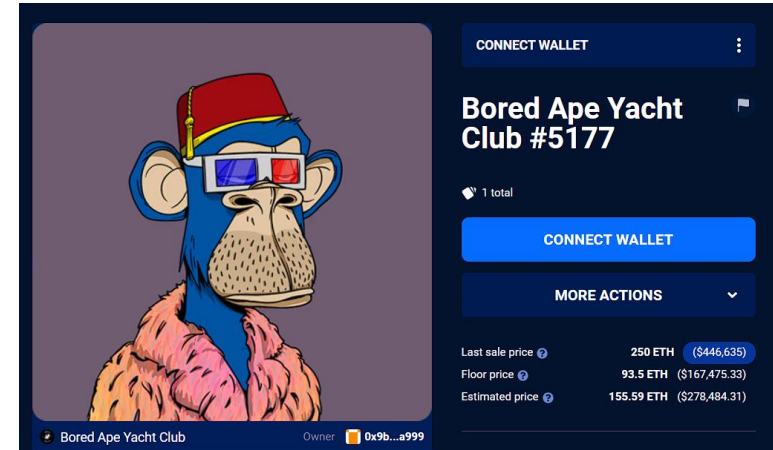


- NFT market is expanding rapidly.
- Pricing and estimation of NFT is not well developed yet because of its subjectivity.
- Our goal is provide effective method to estimate future price of different NFTs.

# Related Works

Study	HR	RSR	VAR	ML	WL
(Kong & Lin, 2021)	*				
(Schaar & Kampakis, 2022)	*				
(Goldberg et al., 2021)	*				
(Nadini et al., 2021)	*	*		*	
(Kireyev & Lin, 2021)	*			*	
(Ante, 2021a,b)			*		
(Dowling, 2022b)					*
(Umar et al., 2022)					*


Most of current analyses are based on Regression



Using historical market data

# Data Overview

- Total 8 collections, 2 of each from Game, Collectible, Art, Metaverse
- From 2017.01.01 to 2022.05.15
- Internal data
  - # of active market wallets
  - # of sales
  - # of unique buyers
  - Total amount of sales in USD
  - Average USD of the collection
- External data
  - NASDAQ
  - ETH/USD
  - USD index



# Linear Regression

$$\hat{y}_t = \hat{\beta}_0 + \hat{\beta}_1 x_{1,t} + \hat{\beta}_2 x_{2,t} + \dots + \hat{\beta}_k x_{k,t}$$

- not genuine forecasts of future value
- just estimate model by linear regression
- can interpret coefficient as feature importance
- feature : {'Number of sales', 'Sales USD', 'Active market wallets', 'Unique buyers'}
- target variable : {'Average USD'}

# Evaluation

- $total\ RMSE = \sqrt{\frac{1}{T} \sum_{t=1}^T (\hat{y}_t - y_t)^2}$  (sklearn.metrics.mean\_squared\_error)


- coefficient of determination (sklearn.metrics.r2\_score)

$$R^2 = \frac{\Sigma(\hat{y}_t - \bar{y})^2}{\Sigma(y_t - \bar{y})^2}$$

# Linear Regression


RMSE: 1716.8893388316112

Coefficient of Determination: 0.15170165831779292



RMSE: 26837.011415394452

Coefficient of Determination: -73390.40486717146




# Feature scaling

- standard feature scaling (sklearn.preprocessing.StandardScaler)
- minmax feature scaling (sklearn.preprocessing.MinMaxScaler)


RMSE: 2920.7048660522487

Coefficient of Determination: -1.4549313883860453



RMSE: 78.55120162980465

Coefficient of Determination: -73390.40486717236





# Lasso Regression (=L1 Regularization)

- automatic feature selection

$$L(w) = \frac{1}{N} \sum_{n=1}^N \frac{1}{2} \left( w^T x^{(n)} - y^{(n)} \right)^2 + \lambda \sum_i^d |w_d|$$

- `sklearn.linear_model.LassoCV`

# Lasso Regression (=L1 Regularization)

Best alpha using built-in LassoCV: 1418431045.719116

Best score using built-in LassoCV: 0.139841

1728.8503341547337

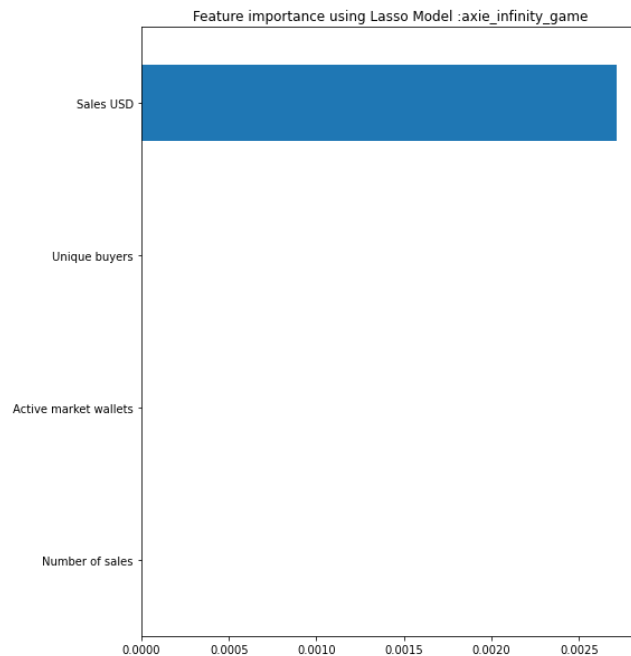

Lasso picked 1 variables and eliminated the other 3 variables

Best alpha using built-in LassoCV: 48899.083275

Best score using built-in LassoCV: -137176.442410

36690.43840677007

Lasso picked 1 variables and eliminated the other 3 variables



# additional feature +Lasso Regression

- add new feature : {'ETH', 'NASDAQ', 'USD'}

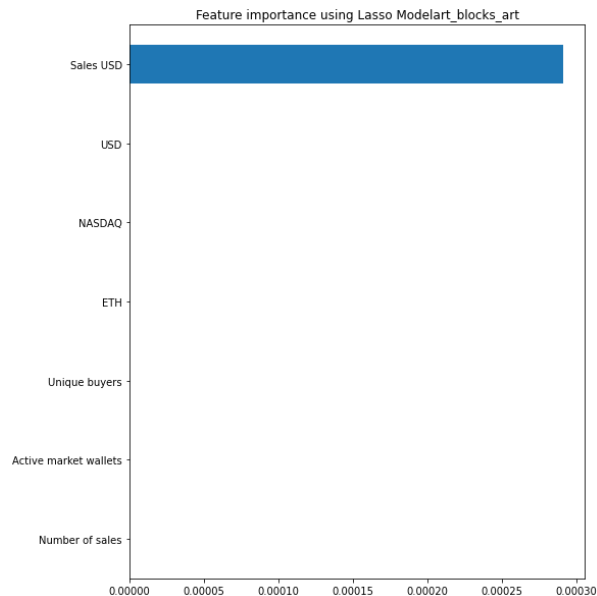
Best alpha using built-in LassoCV: 1420545063.481934

Best score using built-in LassoCV: 0.139417

1729.275958047487

0.13941728794359343

Lasso picked 1 variables and eliminated the other 6 variables



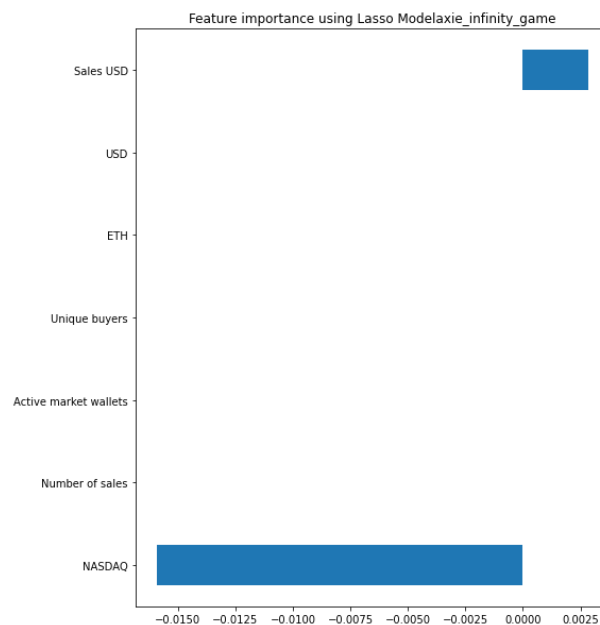
Best alpha using built-in LassoCV: 49035.535694

Best score using built-in LassoCV: -150780.862347

38420.10681044345



-150780.86234651148

Lasso picked 2 variables and eliminated the other 5 variables



# EMA(Exponential Moving Average)

- pandas\_ta library ema function
- EMA10



# Linear regression for EMA10

Best alpha using built-in LassoCV: 678526102.898727



Best score using built-in LassoCV: 0.028498

934.3928054105266

Best alpha using built-in LassoCV: 10753.052293


Best score using built-in LassoCV: -12027.379186

10465.256472109817

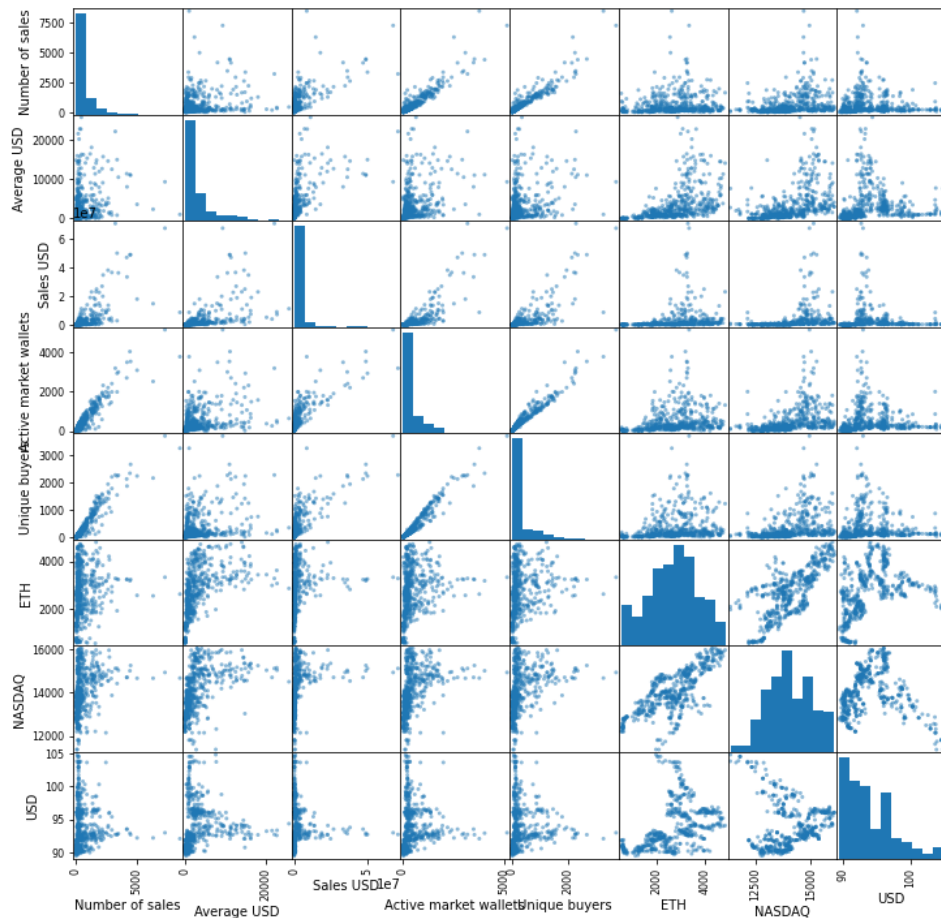


# why?

## correlation between features




art\_blocks\_art

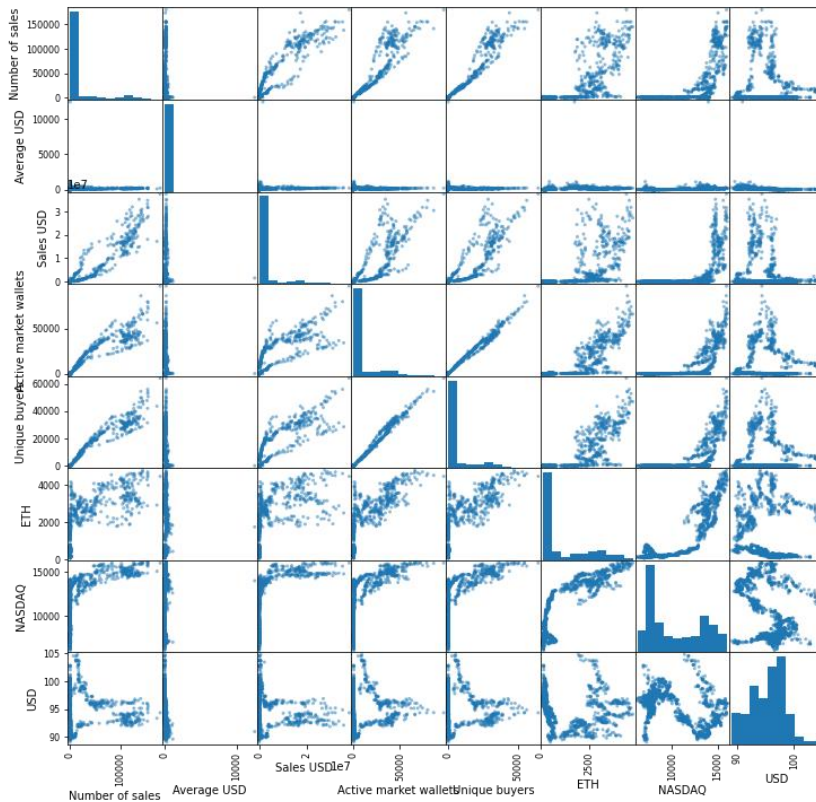


# why?

## correlation between features



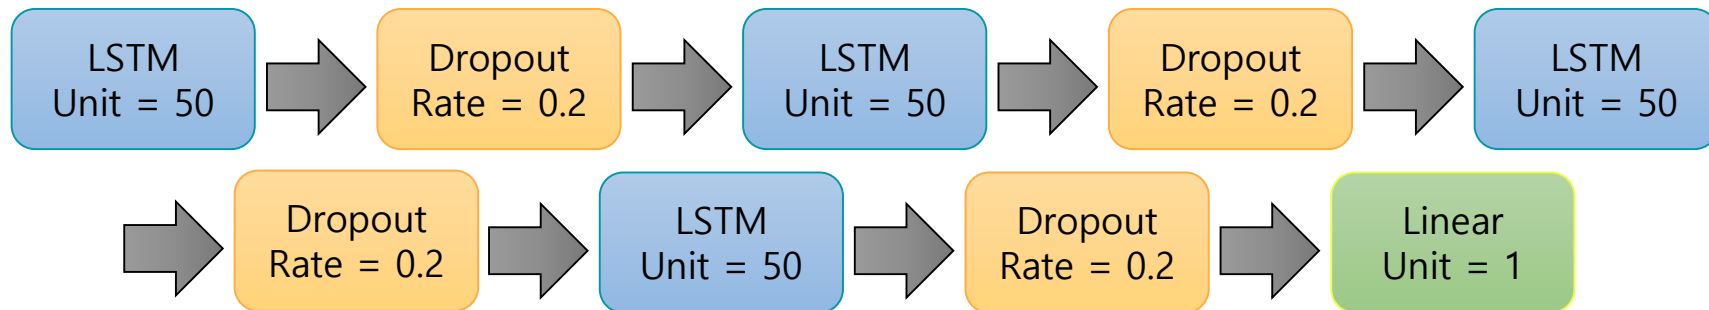
axie\_infinity\_game



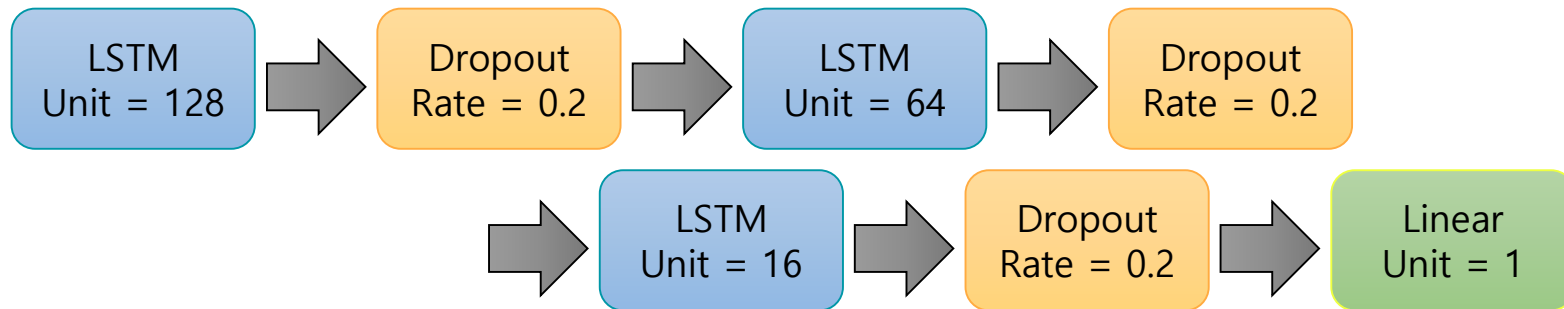
# Stacked LSTM



## Architecture 1






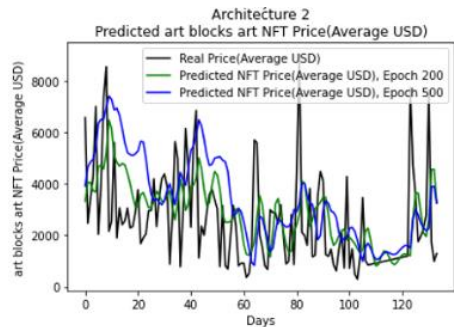



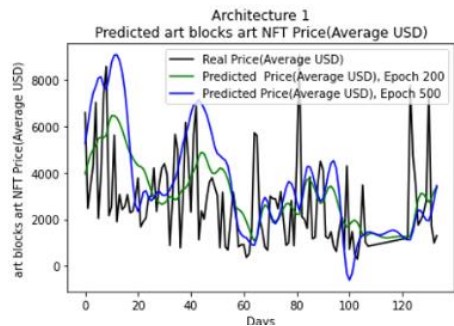
## Architecture 2











Optimizer: ADAM, Loss: Mean squared Error

**Epoch: 200, 500 each**

# Prediction Graph (1)




# Prediction Graph (2)



# RMSE (Based on Real, unscaled value)


Data Category	RSME				Lowest RMSE
	Architecture 1		Architecture 2		
	Epoch 200	Epoch 500	Epoch 200	Epoch 500	
art blocks art	2035	2589	1879	2249	1879
Creepz game	1464	1521	1206	1355	1206
The sandbox metaverse	8112	9279	10262	9284	8112
Superrare art	43569	45175	44324	48966	43569
Bored ape yacht club collection	21582	18245	18262	16691	16691
Axie infinity game	39	42	40	120	39
Crypto punk collection	204144	163752	151422	191073	151422
Decentraland metaverse	9702	9448	7703	9238	7703

# Reinforcement Learning




	Number of sales	Average USD	Sales USD	Active mar
4/25/2021	7	183.62	1285.33	8
4/26/2021	70	198.59	13901.4	18
4/27/2021	46	204.71	9416.47	20
4/28/2021	49	215.45	10557.25	19
4/29/2021	68	219.25	14909.21	30
4/30/2021	297	221.81	65877.04	90
5/1/2021	11159	337.48	3765957.9	140
5/2/2021	1706	2301.68	3926670.5	110
5/3/2021	1142	2657.02	3034317	88
5/4/2021	350	2819.33	986765.68	41
5/5/2021	269	2357.26	634103.49	33
5/6/2021	191	2473.23	472386.45	26
5/7/2021	155	2055.3	318572.1	21
5/8/2021	97	2657.96	257822.49	15
5/9/2021	119	2566.07	305362.86	16
5/10/2021	62	2570.01	159340.91	90
5/11/2021	35	1577.69	55219.31	60
5/12/2021	41	2198.69	90146.39	70
5/13/2021	48	1734.96	83277.97	80
5/14/2021	43	2277.31	97924.37	70


# Reinforcement Learning



# Reinforcement Learning




# Reinforcement Learning







# Reinforcement Learning



# Reinforcement Learning





# Reinforcement Learning



# Policy Network/ Agent

- Inputs an observation, outputs an action (for the agent to take)
- The brain of the agent
- MLP policy: 64 input feature extractor, 64 input FCN



# Observation Space

- Basic NFT data features as the observations
  - 75/25 training/test split
- We also tried adding the following values to the observation:
  - The agent's most recent prediction value
  - Each feature of the most recent observation

# Algorithms

Actor-critic RL algorithms that could support continuous observation and action spaces

- PPO (Proximal Policy Approximation)
- A2C (Advantage Actor Critic)

# Normalization

- In order for Actor critic algorithms to learn well, the observation, reward and especially the actions, should be normalized to reasonable values

10/20/202	93.56	4053.9	15,121.68
10/21/202	93.77	3971.55	15,215.70
10/22/202	93.64	4168.56	15,090.20
10/23/202	93.64	4082.64	15,090.20
10/24/202	93.64	4220.93	15,090.20
10/25/202	93.81	4131.47	15,226.71
10/26/202	93.95	3923.94	15,235.71
10/27/202	93.8	4288.26	15,235.84
10/28/202	93.35	4421.23	15,448.12
10/29/202	94.12	4324.67	15,498.39
10/30/202	94.12	4290.16	15,498.39
10/31/202	94.12	4322.89	15,498.39
11/1/2021	93.88	4595.01	15,595.92



Reward: 2502342



# Normalization of Action

For PPO and A2C, it is critical that the action space be normalized to  $[-1, 1]$

=> Use Tanh as the activation function in the Policy Network to achieve this.



# Normalization of Observation Space

Normalizing the observation was not feasible

=> The agent would be able to predict values only in the normalized environment

=> Not able to predict the actual price values

**Problem:** Action space is limited to  $[-1, 1]$ , whereas the observation is not

**Solution:** Using the action value as a **ratio** to be multiplied to the agent's current location to make a prediction

ex) ratio:  $1 + (0.5 * \text{action})$  => ratio can vary from 0.5 to 1.5!


current agent's location: \$2000, action = 1 => predicted value:  $2000 * 1.5 = 3000$

# Normalization of Reward


- We used the negative absolute difference as initial reward
  - Normalize it to  $[-2,0]$  using moving averages and variances
- => we did not want to give positive rewards in order to get the agent to predict as close to the actual value possible

# Results


- Bored Ape Yacht Club Collection, PPO, 150000 timesteps




< Training Process >



< Entire dataset >




< Test set >




< Smaller set >

# Results


- Bored Ape Yacht Club Collection, A2C, 150000 timesteps




< Training Process >



< Entire dataset >




< Test set >




< Smaller set >

# Results


- Axie Infinity Game, PPO, 150000 timesteps



< Entire dataset >




< Test set >




< Smaller set >

# Results


- Axie Infinity Game, A2C, 150000 timesteps



< Entire dataset >



< Test set >



< Smaller set >

# Evaluation Metrics: RMSE, R2

Bored Ape, PPO  
RMSE: 24523.273  
R2: -0.525

Axie Infinity, PPO  
RMSE: 100.979  
R2: -0.041

Bored Ape, A2C  
RMSE: 19662.734042581218  
R2: 0.11082038653062165

Axie Infinity, A2C  
RMSE: 60.950  
R2: 0.620

# Conclusion



	Linear Regression (Baseline)	Stacked LSTM	Reinforcement Learning
RMSE	26837.011415394	40.8	100.979
R <sup>2</sup> score	-73390.40486717	0.8296477128246	0.62

- LSTM showed the best results => best handles sequential data
- Future work: Combining models - Reinforcement learning model that uses LSTM feature extractors!